

1. See previous submission.
2. See previous submission.
3. In only serial execution, i.e. $p = 1$, the total time taken T_s would be $t\tau$, where τ is the time taken per task. In pipelined execution with arbitrary p , the time taken for the very first task to complete is still τ because the pipeline is cold. However, each subsequent task will then come off the pipeline with a takt time of τ/p . Thus the total time taken would be:

$$T_p = \tau + \frac{(t-1)\tau}{p}$$

The speedup is:

$$\frac{T_s}{T_p} = \frac{t\tau}{\tau + \frac{(t-1)\tau}{p}} = \frac{pt}{p+t+1}$$

4. The minimum serial time is 2.75h. The minimum parallel time is 2.25h.
5. See [Figure 1](#).

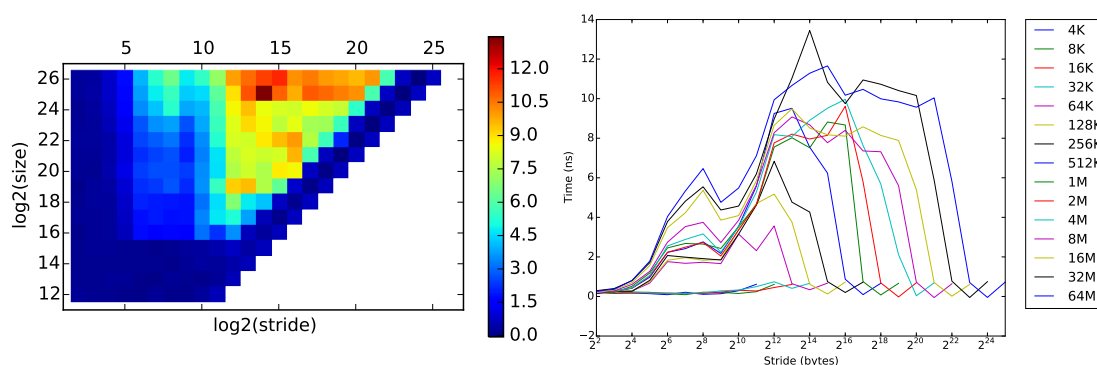


Figure 1: Membench results for the local system

6. See [Figure 2](#).
7. A caveat: when attempting to run `centroid.c` verbatim from the course repo on the cluster, one found that it was consistently giving a timing of 0 for all three functions. The initial hypothesis was that the resolution of the default C timer was insufficient. Following up on Piazza, `centroid.c` was modified to use the OpenMP timer, giving the following results:

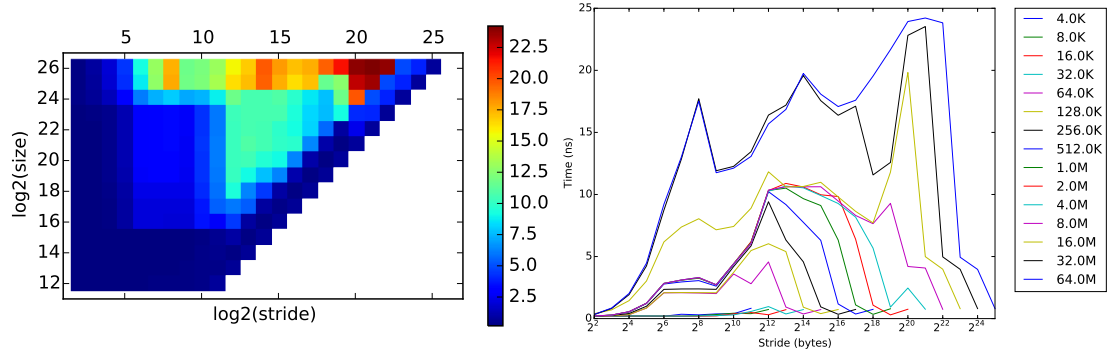


Figure 2: Membench results for the totient cluster

- (a) $4.440892\text{e-}16$ to $1.443290\text{e-}15$
- (b) $4.440892\text{e-}16$ to $1.443290\text{e-}15$
- (c) $1.332268\text{e-}15$ to $1.443290\text{e-}15$

also known as the unfortunate case where the numbers do not really make sense. Theoretically, Implementation B should be the slowest because it does not exploit memory locality or look-ahead, whereas Implementations A and C should be faster because the compiler can optimize for array access from contiguous blocks of memory.