

1. The number of double precision numbers evaluated in the innermost loop is $N+N = 2N$. Each double precision number is 8 bytes, thus the total number of bytes handled in the innermost loop is $16N$. Since this value is larger than the size of the L3 cache, it must be paged in from memory. Thus the memory-based AI of the code is $1/8$.
2. Since the cache is smaller than $8N^2$, it is infeasible to hold the entire matrix multiplication (exactly $8N^2$ bytes) in cache. However, we can easily hold one set of values in cache, say, a row of A for which all columns of B calculated. This improves the arithmetic intensity by a factor of 2, i.e. $1/4$.
3. We start by paging in $16N^2$ bytes accounting for the matrices A and B . The number of operations performed is $2N^3$. This produces C , which has to be transferred back to memory, i.e. another $16N^2$ bytes. Therefore the AI is:

$$\frac{2N^3}{16N^2 + 16N^2} = \frac{N}{16}$$

4. The assumption used here is that C is not held in the cache because the result can be pushed back to memory as it is computed. Thus we need to store $16N^2$ bytes in cache. For the L1 cache, this corresponds to $N = 32$ with an AI of 2. For the L2 cache, this corresponds to $N = 90$ with an AI of 5.625. For the L3 cache, this corresponds to $N = 443$ with an AI of 27.6875.
5. The Flop rate for the machine is $4 * 2 * 8 * 2.4 = 153.6$ GFlops/s (where each FMA instruction is 2 flops). Thus the max AI is $153/25.6 = 5.98$.
6. When $N = 16 * 6 = 96$.
7. Flops increases monotonically with N and asymptotes when the CPU-bound AI is reached.